

Introduction to Probability and Statistics

January 23, 2014




Last time ...

- 1) Formulate a *model* of pairs of sentences.
- 2) *Learn* an instance of the model from *data*.
- 3) Use it to *infer* translations of new inputs.

Why Probability?

- Probability formalizes ...
 - the concept of *models*
 - the concept of *data*
 - the concept of *learning*
 - the concept of *inference* (prediction)



Probability is expectation founded upon partial knowledge.

$$p(x \mid \text{partial knowledge})$$

“Partial knowledge” is an apt description of what we know about language and translation!

Probability Models

- Key components of a probability model
 - The space of events (Ω or S)
 - The assumptions about conditional independence / dependence among events
 - Functions assigning probability (density) to events
 - *We will assume discrete distributions.*

Events and Random Variables

A **random variable** is a function from a random event from a set of possible outcomes (Ω) and a probability distribution (p), a function from outcomes to probabilities.

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$



$$X(\omega) = \omega$$

$$\rho_X(x) = \begin{cases} \frac{1}{6} & \text{if } x = 1, 2, 3, 4, 5, 6 \\ 0 & \text{otherwise} \end{cases}$$

Events and Random Variables

A **random variable** is a function from a random event from a set of possible outcomes (Ω) and a probability distribution (p), a function from outcomes to probabilities.

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$

$$Y(\omega) = \begin{cases} 0 & \text{if } \omega \in \{2, 4, 6\} \\ 1 & \text{otherwise} \end{cases}$$

$$\rho_Y(y) = \begin{cases} \frac{1}{2} & \text{if } y = 0, 1 \\ 0 & \text{otherwise} \end{cases}$$



What is our event space?

**What are our random
variables?**

Probability Distributions

A probability distribution (p_X) assigns probabilities to the values of a random variable (X).

There are a couple of philosophically different ways to define probabilities, but we will give only the invariants in terms of **random variables**.

$$\sum_{x \in \mathcal{X}} \rho_X(x) = 1$$

$$\rho_X(x) \geq 0 \quad \forall x \in \mathcal{X}$$


Probability distributions of a random variable may be specified in a number of ways.

Specifying Distributions

- Engineering/mathematical convenience
- Important techniques in this course
 - Probability mass functions
 - Tables (“stupid multinomials”)
 - Log-linear parameterizations (maximum entropy, random field, multinomial logistic regression)
 - Construct random variables from other r.v.’s with known distributions

Sampling Notation

$$x = 4 \times z + 1.7$$


Variable

Expression

Sampling Notation

$$x = 4 \times z + 1.7$$

$$y \sim \text{Distribution}(\theta)$$

Distribution

Random variable

Parameter

Sampling Notation

$$y \sim \text{Distribution}(\boldsymbol{\theta})$$

$$y' = y \times x$$

Multivariate r.v.'s

Probability theory is particularly useful because it lets us reason about (cor)related and dependent events.

A joint probability distribution is a probability distribution over r.v.'s with the following form:

$$Z = \begin{bmatrix} X(\omega) \\ Y(\omega) \end{bmatrix}$$

$$\sum_{x \in \mathcal{X}, y \in \mathcal{Y}} \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) = 1 \quad \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) \geq 0 \quad \forall x \in \mathcal{X}, y \in \mathcal{Y}$$

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$



$$X(\omega) = \omega$$

$$\begin{aligned} \Omega = \{ & (1, 1), (1, 2), (1, 3), (1, 4), (1, 5), (1, 6), \\ & (2, 1), (2, 2), (2, 3), (2, 4), (2, 5), (2, 6), \\ & (3, 1), (3, 2), (3, 3), (3, 4), (3, 5), (3, 6), \\ & (4, 1), (4, 2), (4, 3), (4, 4), (4, 5), (4, 6), \\ & (5, 1), (5, 2), (5, 3), (5, 4), (5, 5), (5, 6), \\ & (6, 1), (6, 2), (6, 3), (6, 4), (6, 5), (6, 6), \} \end{aligned}$$



$$X(\omega) = \omega_1 \quad Y(\omega) = \omega_2$$

$$\rho_{X,Y}(x, y) = \begin{cases} \frac{1}{36} & \text{if } (x, y) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$



$$X(\omega) = \omega$$

$$X(\omega) = \omega_1 \quad Y(\omega) = \omega_2$$

$$\rho_{X,Y}(x,y) = \begin{cases} \frac{1}{36} & \text{if } (x,y) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$



$$X(\omega) = \omega$$

$$\begin{aligned} \Omega = \{ & (1, 1), (1, 2), (1, 3), (1, 4), (1, 5), (1, 6), \\ & (2, 1), (2, 2), (2, 3), (2, 4), (2, 5), (2, 6), \\ & (3, 1), (3, 2), (3, 3), (3, 4), (3, 5), (3, 6), \\ & (4, 1), (4, 2), (4, 3), (4, 4), (4, 5), (4, 6), \\ & (5, 1), (5, 2), (5, 3), (5, 4), (5, 5), (5, 6), \\ & (6, 1), (6, 2), (6, 3), (6, 4), (6, 5), (6, 6), \} \end{aligned}$$



$$X(\omega) = \omega_1 \quad Y(\omega) = \omega_2$$

$$\rho_{X,Y}(x, y) = \begin{cases} \frac{x+y}{252} & \text{if } (x, y) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

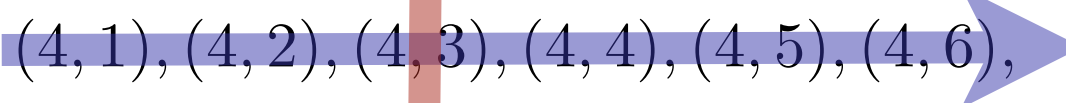
Marginal Probability

$$p(X = x, Y = y) = \rho_X(x, y)$$

$$p(X = x) = \sum_{y' \in \mathcal{Y}} p(X = x, Y = y')$$

$$p(Y = y) = \sum_{x' \in \mathcal{X}} p(X = x', Y = y)$$

$\Omega = \{(1, 1), (1, 2), (1, 3), (1, 4), (1, 5), (1, 6),$
 $(2, 1), (2, 2), (2, 3), (2, 4), (2, 5), (2, 6),$
 $(3, 1), (3, 2), (3, 3), (3, 4), (3, 5), (3, 6),$
 $(4, 1), (4, 2), (4, 3), (4, 4), (4, 5), (4, 6),$
 $(5, 1), (5, 2), (5, 3), (5, 4), (5, 5), (5, 6),$
 $(6, 1), (6, 2), (6, 3), (6, 4), (6, 5), (6, 6), \}$


$$p(X = 4) = \sum_{y' \in [1, 6]} p(X = 4, Y = y')$$


$$p(Y = 3) = \sum_{x' \in [1, 6]} p(X = x', Y = 3)$$

$$\rho_{X,Y}(x,y) = \begin{cases} \frac{1}{36} & \text{if } (x,y) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

$$\Omega = \{(1,1), (1,2), (1,3), (1,4), (1,5), (1,6), \\ (2,1), (2,2), (2,3), (2,4), (2,5), (2,6), \\ (3,1), (3,2), (3,3), (3,4), (3,5), (3,6), \\ (4,1), (4,2), (4,3), (4,4), (4,5), (4,6), \\ (5,1), (5,2), (5,3), (5,4), (5,5), (5,6), \\ (6,1), (6,2), (6,3), (6,4), (6,5), (6,6), \}$$

$$\frac{6}{36} = \frac{1}{6}$$

$$\rho_{X,Y}(x,y) = \begin{cases} \frac{x+y}{252} & \text{if } (x,y) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

$$\Omega = \{(1,1), (1,2), (1,3), (1,4), (1,5), (1,6), \\ (2,1), (2,2), (2,3), (2,4), (2,5), (2,6), \\ (3,1), (3,2), (3,3), (3,4), (3,5), (3,6), \\ (4,1), (4,2), (4,3), (4,4), (4,5), (4,6), \\ (5,1), (5,2), (5,3), (5,4), (5,5), (5,6), \\ (6,1), (6,2), (6,3), (6,4), (6,5), (6,6), \}$$

$$\frac{4+1+4+2+4+3+4+4+4+5+4+6}{252} = \frac{45}{252}$$

Conditional Probability

The conditional probability of one random variable given another is defined as follows:

$$p(X = x \mid Y = y) = \frac{p(X = x, Y = y)}{p(Y = y)} = \frac{\text{joint probability}}{\text{marginal}}$$

Given that $p(y) \neq 0$

Conditional probability distributions are useful for specifying joint distributions since:

$$p(x \mid y)p(y) = p(x, y) = p(y \mid x)p(x)$$

Why might this be useful?

Conditional Probability Distributions

A conditional probability distribution is a probability distribution over r.v.'s X and Y with the form $\rho_{X|Y=y}(x)$.

$$\sum_{x \in \mathcal{X}} \rho_{X|Y=y}(x) = 1 \quad \forall y \in \mathcal{Y}$$

Chain rule

The **chain rule** is derived from a repeated application of the definition of conditional probability:

$$p(a, b, c, d)$$



Use as many times as necessary!

Bayes' Rule

Posterior Likelihood Prior

$p(x | y) = \frac{p(y | x)p(x)}{p(y)} \left(= \frac{p(y | x)p(x)}{\sum_{x'} p(y | x')p(x')} \right)$

Evidence

The diagram illustrates the components of Bayes' Rule. The word 'Posterior' has an arrow pointing to $p(x | y)$. The word 'Likelihood' has an arrow pointing to $p(y | x)$. The word 'Prior' has an arrow pointing to $p(x)$. The word 'Evidence' has an arrow pointing to $p(y)$ in the denominator. The equation is presented in two forms: a simplified version and a version with a summation over all possible x' .

Independence

Two random variables are independent iff

$$p(X = x, Y = y) = p(X = x)p(Y = y)$$

Equivalently, (use def. of cond. prob to prove)

$$p(X = x \mid Y = y) = p(X = x)$$

Equivalently again:

$$p(Y = y \mid X = x) = p(Y = y)$$

“Knowing about X doesn’t tell me about Y ”

$$\rho_{X,Y}(x,y) = \begin{cases} \frac{1}{36} & \text{if } (x,y) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

$$\Omega = \{(1,1), (1,2), (1,3), (1,4), (1,5), (1,6), \\ (2,1), (2,2), (2,3), (2,4), (2,5), (2,6), \\ (3,1), (3,2), (3,3), (3,4), (3,5), (3,6), \\ (4,1), (4,2), (4,3), (4,4), (4,5), (4,6), \\ (5,1), (5,2), (5,3), (5,4), (5,5), (5,6), \\ (6,1), (6,2), (6,3), (6,4), (6,5), (6,6), \}$$

$$\rho_{X,Y}(x,y) = \begin{cases} \frac{x+y}{252} & \text{if } (x,y) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

$$\Omega = \{(1,1), (1,2), (1,3), (1,4), (1,5), (1,6), \\ (2,1), (2,2), (2,3), (2,4), (2,5), (2,6), \\ (3,1), (3,2), (3,3), (3,4), (3,5), (3,6), \\ (4,1), (4,2), (4,3), (4,4), (4,5), (4,6), \\ (5,1), (5,2), (5,3), (5,4), (5,5), (5,6), \\ (6,1), (6,2), (6,3), (6,4), (6,5), (6,6), \}$$

Independence

Independence has **practical benefits**. Think about how many parameters you need for a naive parameterization of $\rho_{X,Y}(x,y)$ **vs** $\rho_X(x)$ and $\rho_Y(y)$

$$O(xy) \text{ vs } O(x + y)$$

Conditional Independence

Two equivalent statements of conditional independence:

$$p(a, c \mid b) = p(a \mid b)p(c \mid b)$$

and:

$$p(a \mid b, c) = p(a \mid b)$$

“If I know B , then C doesn’t tell me about A ”

Conditional Independence

$$\begin{aligned} p(a, b, c) &= p(a \mid b, c)p(b, c) \\ &= p(a \mid b, c)p(b \mid c)p(c) \end{aligned}$$

“If I know B , then C doesn’t tell me about A ”

$$p(a \mid b, c) = p(a \mid b)$$

$$\begin{aligned} p(a, b, c) &= p(a \mid b, c)p(b, c) \\ &= p(a \mid b, \textcolor{red}{c})p(b \mid c)p(c) \\ &= p(a \mid b)p(b \mid c)p(c) \end{aligned}$$

Do we need more parameters or fewer parameters in conditional independence?

Independence

- Some variables are independent In Nature
 - How do we know?
- Some variables we *pretend* are independent for computational convenience
 - Examples?
- Assuming independence is equivalent to letting our model “forget” something that happened in its past
 - What should we forget in language?

A Word About Data

- When we formulate our models there will be two kinds of random variables: observed and latent
- Observed: words, sentences(?), parallel corpora, web pages, formatting...
- Latent: parameters, syntax, “meaning”, word alignments, translation dictionaries...

Interlingua
“meaning”

```
report_event[  
  factivity=true  
  explode(e, bomb, car)  
  loc(e, downtown)  
]
```

```
explodieren  
:arg0 Bomb  
:arg1 car  
:loc Innenstadt  
:tempus imperf
```

```
detonate  
:arg0 bomb  
:arg1 car  
:loc downtown  
:time past
```

In der Innenstadt explodierte eine Autobombe

A car bomb exploded downtown

In der Innenstadt explodierte eine Autobombe

A car bomb exploded downtown

Hidden

Observed

Garcia and associates .

Garcia y asociados .

Carlos Garcia has three associates .

Carlos Garcia tiene tres asociados .

his associates are not strong .

sus asociados no son fuertes .

Garcia has a company also .

Garcia tambien tiene una empresa .

its clients are angry .

sus clientes estan enfadados .

the associates are also angry .

los asociados tambien estan enfadados .

the clients and the associates are enemies .

los clientes y los asociados son enemigos .

the company has three groups .

la empresa tiene tres grupos .

its groups are in Europe .

sus grupos estan en Europa .

the modern groups sell strong pharmaceuticals .

los grupos modernos venden medicinas fuertes .

the groups do not sell zanzanine .

los grupos no venden zanzanina .

the small groups are not modern .

los grupos pequenos no son modernos .

Hidden

Garcia and associates .

\ \ /
Garcia y asociados .

Carlos Garcia has three associates .

\ \ | | /
Carlos Garcia tiene tres asociados .

his associates are not strong .

| \ X /
sus asociados no son fuertes .

Garcia has a company also .

| \ X X /
Garcia tambien tiene una empresa .

its clients are angry .

/ / | \
sus clientes estan enfadados .

the associates are also angry .

/ / X \
los asociados tambien estan enfadados .

the clients and the associates are enemies .

\ \ | / / / /
los clientes y los asociados son enemigos .

the company has three groups .

\ | / / /
la empresa tiene tres grupos .

its groups are in Europe .

/ | | \
sus grupos estan en Europa .

the modern groups sell strong pharmaceuticals .

| X \ X
los grupos modernos venden medicinas fuertes .

the groups do not sell zanzanine .

| | / / /
los grupos no venden zanzanina .

the small groups are not modern .

/ X X \
los grupos pequenos no son modernos .

Learning

- Let's say we have formulated a model of a phenomenon
 - Made independence assumptions
 - Figured out what kinds of parameters we want
- Let's say we have collected data we assume to be generated by this model
 - E.g. some parallel data

What do we do now?

Parameter Estimation

- Inputs
 - Given a model with unspecified parameters
 - Given some data
- Goal: learn model parameters
- How?
 - Find parameters that make the model make predictions that look like the data do
 - What do we mean “look like the data?”
 - Probability (other options: accuracy, moment matching)

Strategies

- **Maximum likelihood estimation**
 - What is the *probability* of generating the data?
- **Accuracy**
 - Using an auxiliary similarity function, find parameters that maximize the (expected?) accuracy of data
- **Bayesian techniques**



$p(\text{heads})$





$1 - p(\text{heads})$

$p(\text{heads})$?



$$p(\text{data}) = p(\text{heads})^7 \times [1 - p(\text{heads})]^3$$





Optimization


- For the most part, we will be working with maximum likelihood estimation
- The general recipe is:
 - Come up with an expression of the likelihood of your probability model, as a function of **data** and the model **parameters**
 - Set the parameters to maximize the likelihood
 - This optimization is generally difficult
 - You must respect any constraints on the parameters (>0 , sum to 1, etc)
 - There may not be analytical solutions (log-linear models)

Probability lets us

- 1) Formulate a *model* of pairs of sentences.
- 2) *Learn* an instance of the model from *data*.
- 3) Use it to *infer* translations of new inputs.

Supplemental Reading

- If this was unfamiliar to you, then please read Chapter 3 from the textbook "Statistical Machine Translation" by Philipp Koehn



Announcements

- If you haven't done so already, complete HW 0 by today at 11:59pm.
- Office hours are set. Jonny: Wednesdays at 2pm (Levine 5th floor bump space), Chris: Mondays at 10:30am (Levine 506)
- HW1 will be released this weekend, and due on Thursday Feb 6. I strongly encourage you to do it before the end of the course selection period (Feb 3).

Announcements

- Grading has been set
- 7 homework assignments, 10 points each
- 1 in-class presentation, 10 points
- 1 final project with writeup and code 20 points
- I'll post a description of the requirements on the web page, and then send a note to piazza